



## Region of interest pooling explained

February 28, 2017 / 42 Comments / in Data science, Deep learning, Machine learning / by Tomasz Grel

**Region of interest pooling (also known as RoI pooling) is an operation widely used in object detection tasks using convolutional neural networks. For example, to detect multiple cars and pedestrians in a single image. Its purpose is to perform max pooling on inputs of nonuniform sizes to obtain fixed-size feature maps (e.g. 7×7).**

We've just released an open-source implementation of RoI pooling layer for TensorFlow (you can find it [here](#)). In this post, we're going to say a few words about this interesting neural network layer. But first, let's start with some background.

Two major tasks in computer vision are object classification and object detection. In the first case the system is supposed to correctly label the dominant object in an image. In the second case it should provide correct labels and locations for all objects in an image. Of course there are other

### POPULAR POSTS

**Don't waste the power. AI supply chain management**

December 13, 2018 - 4:09 pm

**Powering up demand forecasting with machine learning**

December 6, 2018 - 2:41 pm

**AI Monthly digest #3: Artificial Intelligence in science getting big**

December 6, 2018 - 2:41 pm

### CATEGORIES

Big data & Spark

Data science

Deep learning

Machine learning

Neptune

Reinforcement learning

Seahorse

### FOLLOW OUR NEWSLETTER

Subscribe

You can modify your privacy settings and unsubscribe from our lists at any time (see our [privacy policy](#)).

bounding boxes around any object from a previously specified set of categories and assign a class to each of them. For example, let's say we're developing an algorithm for self-driving cars and we'd like to use a camera to detect other cars, pedestrians, cyclists, etc. — our dataset might look like [this](#).

In this case we'd have to draw a box around every significant object and assign a class to it. This task is more challenging than classification tasks such as [MNIST](#) or [CIFAR](#). On each frame of the video, there might be multiple objects, some of them overlapping, some poorly visible or occluded. Moreover, for such an algorithm, performance can be a key issue. In particular for autonomous driving we have to process tens of frames per second.

So how do we solve this problem?

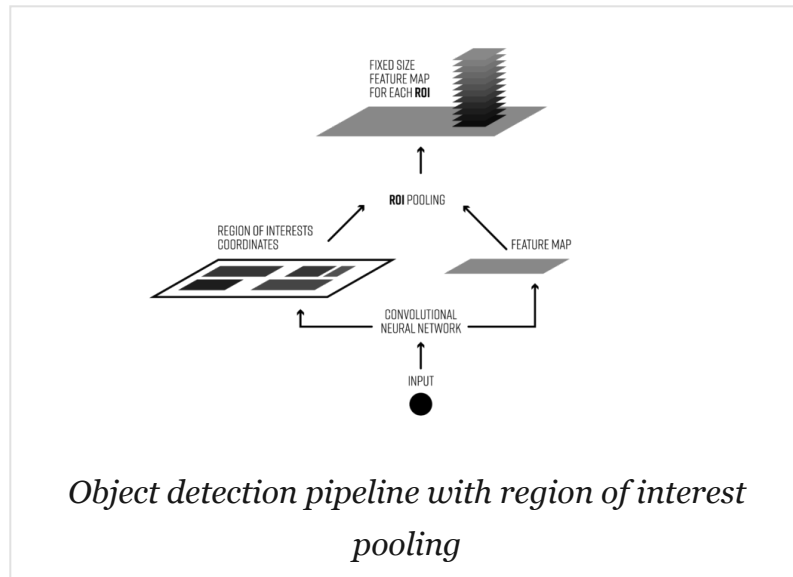
**Related: [Playing Atari with deep reinforcement learning - deepsense.ai's approach](#)**

## Typical architecture

The object detection architecture we're going to be talking about today is broken down in two stages:

1. Region proposal: Given an input image find all possible places where objects can be located. The output of this stage should be a list of bounding boxes of likely positions of objects. These are often called region proposals or regions of interest. There are quite a few methods for this task, but we're not going to talk about them in this post.

background. Here we could use a deep convolutional network.



Usually in the proposal phase we have to generate a lot of regions of interest. Why? If an object is not detected during the first stage (region proposal), there's no way to correctly classify it in the second phase. That's why it's extremely important for the region proposals to have a high recall. And that's achieved by generating very large numbers of proposals (e.g., a few thousands per frame). Most of them will be classified as background in the second stage of the detection algorithm.

Some problems with this architecture are:

- Generating a large number of regions of interest can lead to performance problems. This would make real-time object detection difficult to implement.

train all the components of the system in one run (which would yield much better results)

That's where region of interest pooling comes into play.

**Related: Optimize Spark with  
DISTRIBUTE BY & CLUSTER BY**

## Region of interest pooling — description

Region of interest pooling is a neural-net layer used for object detection tasks. It was first proposed by Ross Girshick in April 2015 (the article can be found [here](#)) and it achieves a significant speedup of both training and testing. It also maintains a high detection accuracy. The layer takes two inputs:

1. A fixed-size feature map obtained from a deep convolutional network with several convolutions and max pooling layers.
2. An  $N \times 5$  matrix of representing a list of regions of interest, where  $N$  is a number of RoIs. The first column represents the image index and the remaining four are the coordinates of the top left and bottom right corners of the region.



An image from the Pascal VOC dataset annotated with region proposals (the pink rectangles)

What does the RoI pooling actually do? For every region of interest from the input list, it takes a section of the input feature map that corresponds to it and scales it to some pre-defined size (e.g.,  $7 \times 7$ ).

The scaling is done by:

1. Dividing the region proposal into equal-sized sections (the number of which is the same as the dimension of the output)
2. Finding the largest value in each section
3. Copying these max values to the output buffer

The result is that from a list of rectangles with different sizes we can quickly get a list of corresponding feature maps with a fixed size. Note that the dimension of the RoI pooling output doesn't actually depend on the size of the input feature map nor on the size of the region proposals. It's determined solely by the number of sections we divide the proposal into. What's the benefit of RoI

same input feature map for all of them. Since computing the convolutions at early stages of processing is very expensive, this approach can save us a lot of time.

[Related: Deep learning for satellite imagery via image segmentation](#)


## Region of interest pooling — example

Let's consider a small example to see how it works. We're going to perform region of interest pooling on a single 8×8 feature map, one region of interest and an output size of 2×2. Our input feature map looks like this:

input


0.88	0.44	0.14	0.16	0.37	0.77	0.96	0.27
0.19	0.45	0.57	0.16	0.63	0.29	0.71	0.70
0.66	0.26	0.82	0.64	0.54	0.73	0.59	0.26
0.85	0.34	0.76	0.84	0.29	0.75	0.62	0.25
0.32	0.74	0.21	0.39	0.34	0.03	0.33	0.48
0.20	0.14	0.16	0.13	0.73	0.65	0.96	0.32
0.19	0.69	0.09	0.86	0.88	0.07	0.01	0.48
0.83	0.24	0.97	0.04	0.24	0.35	0.50	0.91

Let's say we also have a region proposal (top left, bottom right coordinates): (0, 3), (7, 8). In the picture it would look like this:



Normally, there'd be multiple feature maps and multiple proposals for each of them, but we're keeping things simple for the example.


By dividing it into (2×2) sections (because the output size is 2×2) we get:




Notice that the size of the region of interest doesn't have to be perfectly divisible by the number of pooling sections (in this case our RoI is 7×5 and we have 2×2 pooling sections).

The max values in each of the sections are:





And that's the output from the Region of Interest pooling layer. Here's our example presented in form of a nice animation:



What are the most important things to remember about RoI Pooling?

- It's used for object detection tasks
- It allows us to reuse the feature map from the convolutional network
- It can significantly speed up both train and test time
- It allows to train object detection systems in an end-to-end manner



examples on how to use region of interest pooling with Neptune and TensorFlow.

## References:

- Girshick, Ross. "Fast r-cnn." Proceedings of the IEEE International Conference on Computer Vision. 2015.
- Girshick, Ross, et al. "Rich feature hierarchies for accurate object detection and semantic segmentation." Proceedings of the IEEE conference on computer vision and pattern recognition. 2014.
- Sermanet, Pierre, et al. "Overfeat: Integrated recognition, localization and detection using convolutional networks." arXiv preprint arXiv:1312.6229. 2013.

## Related Posts

**Crime forecasting – ‘Minority Report’ realized**

**Don't waste the power. AI supply chain management**

Share this entry



42

REPLIES



**Jere**

March 1, 2017 at 12:22 pm

Quite informative. Waiting for your next post

[Reply](#)



**Rui**

March 31, 2017 at 3:26 pm

typo: "Let's say we also have a region proposal (top left, bottom right coordinates): (5,0), (10,7)", should be (0, 7)

[Reply](#)



**Filip Novotny**

April 15, 2017 at 4:35 pm

Hello,

in this case we'd have to "draw a box around every significant object" and assign a class to it. This task is more challenging [...]

So how do we solve this problem?

From what I understand, you say that Fast(er) R-CNN allow us to predict locations of objects without any ground truth boxes: just run the thing and candidates class-agnostic boxes will pop out.

However, this is a quote from the Fast r-cnn paper:

"Third, the network is modified to take two data inputs: a list of images and a \*list of Rols\* in those images."

There seems to be a need for ground truth boxes. Am I missing something or is there a mistake in the article?

Thanks,

F.

[Reply](#)



**Tomasz Grel**

April 24, 2017 at 8:22 am

Hello,

You definitely need ground truth boxes to train the model and at test time, the model is

every significant object" I meant that the boxes are generated at test time, at train time you have to provide the ground truth boxes and labels.

[Reply](#)



**Jamie**

April 21, 2017 at 2:17 am

Many thanks for the condensed explanation 😊

[Reply](#)



**Krupa**

June 2, 2017 at 7:46 am

I think I'm missing something here. How is the mapping from the input ROI to the feature map ROI done?

[Reply](#)



**Tomasz Grel**

June 2, 2017 at 1:04 pm

Usually the coordinates resulting from the Region Proposal Network are not compatible with the input feature map to the ROI pooling. This is mitigated by careful post processing of those coordinates.

working example here:

<https://blog.deepsense.ai/region-of-interest-pooling-in-tensorflow-example/>

[Reply](#)



**Alex Wang**

August 24, 2017 at 12:16 pm

Hi:

I am confused by this:

‘For every region of interest from the input list, it takes a section of the input feature map that corresponds to it’

for example:

a image of size (299,299), and a region of interest (cat , 50,40, 30, 30), how can I find the region in feature map corresponding to this ROI?

Thanks.

[Reply](#)



**Tomasz Grel**

August 28, 2017 at 8:17 am

In our example we used the ROI format of top left and bottom right corner of the rectangle, so the ROI (50,40,30,30) makes little sense (it should rather be (30,30,50,40) The corresponding ROI is the section of the feature

Reply



**Ilp**

October 26, 2017 at 6:43 am

The ground truth boxes are based on images, but the rois are based on feature maps, how did they correspond to each other? I'm confused about this. Thank you.

Reply



**Tomasz Grel**

October 26, 2017 at 8:11 am

In general you might need to postprocess the boxes so that they match the feature maps. If you need a reference you can find an example along with TensorFlow code in one of our other posts:

<https://blog.deepsense.ai/region-of-interest-pooling-in-tensorflow-example/>

Reply



**Krzysztof**

July 20, 2018 at 2:27 pm

As far as I understand their code, they are using VGG without last pooling layer, so they transform image (224 x 224

dividing them by 16.

I have the feeling that original authors did the same, as this step isn't explained in the paper and dividing by some factor is most obvious way of doing it.

[Reply](#)



**Sai**

December 8, 2017 at 4:29 pm

Great article. Can you elaborate a bit in why the pooling sections the way they are? Is there any intuition for it

[Reply](#)



**Tomasz Grel**

December 8, 2017 at 4:32 pm

The intuition is that you want to partition the image in roughly equal parts, but the size of the output is fixed and the input image can have any size.

[Reply](#)



**Sai**

December 8, 2017 at 10:44 pm

Thanks for the reply! My confusion is about different possible



rectangles we choose  
right?

[Reply](#)



**Sai**

December 8, 2017 at 10:44 pm

Thanks for the timely  
reply! My confusion is  
about different possible  
permutations of 4 pooling  
rectangles. The output  
will differ based on the 4  
rectangles we choose  
right?

[Reply](#)



**Tomasz Grel**

December 11, 2017 at 8:03  
am

Yes, you're right it  
would be different,  
but we don't really  
care about it. It's  
perfectly sufficient  
to just pick one  
deterministic  
method and stick  
with it.

[Reply](#)

confusion for beginners.

“Let’s say we also have a region proposal (top left, bottom right coordinates): (0, 3), (7, 8). In the picture it would look like this:”  
if 1-index is used, top left coordinate would be (1,3) instead of (0,3)  
according to the picture used

[Reply](#)



**Ryan**

December 21, 2017 at 3:26 am

I think there’s a typo in case there’s confusion for beginners.  
“Let’s say we also have a region proposal (top left, bottom right coordinates): (0, 3), (7, 8). In the picture it would look like this:”  
if 1-index is used, top left coordinate would be (1,3) instead of (0,3)  
according to the picture used

[Reply](#)



**sunya**

February 2, 2018 at 12:03 pm

what happens if the rois get too small, say (1\*1) in the feature map?

[Reply](#)

**sunya**

/home/sunya1989/mi/lib/python3.5/s  
packages/roi\_pooling/roi\_pooling.so:  
undefined symbol:  
\_ZTIN10tensorflow8OpKernelE'

[Reply](#)



**Tomasz Grel**

February 5, 2018 at 9:03 am

This looks like a low-level C++ issue. Please check that you're using the right gcc version and the right TensorFlow version.

[Reply](#)



**Tomasz Grel**

February 5, 2018 at 9:00 am

I don't think I've tested this scenario so I'm not sure. If you test this and notice incorrect behavior please create an issue on our github page.

[Reply](#)



**AVi**

March 18, 2018 at 11:00 am

did you understand that, what will happen if rois get too small? for example if we want output from ro pooling as 2x2 but our



**sunya**

February 2, 2018 at 12:03 pm

what happens if the rois get too small, say (1\*1) in the feature map?

[Reply](#)



**sunya**

February 2, 2018 at 12:18 pm

does it perform rescale? (transform rois to the feature map)

[Reply](#)



**Tomasz Grel**

February 5, 2018 at 9:02 am

It just performs the region of interest pooling, nothing else.

[Reply](#)



**sunya**

February 2, 2018 at 12:18 pm

does it perform rescale? (transform rois to the feature map)

[Reply](#)

**sunya**

February 3, 2018 at 1:48 pm

packages/roi\_pooling/roi\_pooling.so:  
undefined symbol:  
\_ZTIN10tensorflow8OpKernelE'

[Reply](#)



**Tomasz Grel**

February 5, 2018 at 9:02 am

It just performs the  
region of interest pooling,  
nothing else.

[Reply](#)



**Tomasz Grel**

February 5, 2018 at 9:03 am

This looks like a low-level  
C issue. Please check that  
you're using the right gcc  
version and the right  
TensorFlow version.

[Reply](#)



**Tomasz Grel**

February 5, 2018 at 9:00 am

I don't think I've tested this  
scenario so I'm not sure. If you  
test this and notice incorrect  
behavior please create an issue  
on our github page.

[Reply](#)

will happen if rois get too small?  
for example if we want output  
from ro pooling as  $2 \times 2$  but our  
roi get smaller than that ( $1 \times 1$ ).  
Do we neglect this roi?

[Reply](#)



**lufficc**

March 7, 2018 at 8:57 am

what if region proposal size is  
smaller than fixed-size feature maps  
size??

[Reply](#)



**Tomasz Grel**

March 12, 2018 at 11:45 am

This situation is perfectly  
normal. Usually proposals will  
cover only a small patch of the  
image where the interesting  
object is located. The purpose of  
the RoI Pooling layer is to cut  
this interesting patch out of the  
feature map and feed it to the  
rest of the network for further  
classification.

[Reply](#)

smaller than fixed-size feature maps  
size??

[Reply](#)



**Tomasz Grel**

March 12, 2018 at 11:45 am

This situation is perfectly normal. Usually proposals will cover only a small patch of the image where the interesting object is located. The purpose of the RoI Pooling layer is to cut this interesting patch out of the feature map and feed it to the rest of the network for further classification.

[Reply](#)



**Neeraj Sajjan**

June 25, 2018 at 11:37 am

In the code of the original implementation[[https://github.com/rbgirs/rcnn/blob/master/matlab/fast\\_rcnn\\_im\\_de](https://github.com/rbgirs/rcnn/blob/master/matlab/fast_rcnn_im_de) and many other implementations[eg: [https://github.com/yhenon/keras-frcnn/blob/master/keras\\_frcnn/RoiPooling](https://github.com/yhenon/keras-frcnn/blob/master/keras_frcnn/RoiPooling) the approach to roi pooling is different from what you have presented here.

In the implementation presented here,



floor( $\lceil 7/2, 5/2 \rceil$ ), i.e., [3,2]. Rest of the boxes will have their sizes adjusted to be [4,2], [3,3] and [4,3].

But according to the original implementation and in several others, they use the same [3,2] as size for all boxes. This leads to ignoring some of the activations in roi proposal.

Am i correct in my above observation and if so, shouldnt your implementation be better than the original?

[Reply](#)



**Tomasz Grel**

June 29, 2018 at 8:29 am

I would indeed assume that ignoring parts of the feature map activations may result in lower accuracy.

[Reply](#)



**Krzysztof**

July 20, 2018 at 8:44 am

There is one part of RoI pooling, which isn't clear for me – you've wrote:

“What does the RoI pooling actually do? For every region of interest from

How do we map coordinates of RoI on original image to coordinates in last layer of ConvNet? For the nets proposed in paper (AlexNet, VGG) the problem is how to map coordinates from (224, 224, 3) to (7, 7, 512). It's not trivial at all. Even if we omit last pooling layer in convNet, then it's still (14, 14, 256).

[Reply](#)



**Tomasz Grel**

July 20, 2018 at 8:52 am

Great question. This is not trivial, but no matter what backbone you're using (VGG, Resnet, FPP) you should be able to postprocess the RoI coordinates to map them into original image coordinates. Most papers treat this step as a technical detail so they don't really talk about it. If you really want to dive deep into this my advice would be to read the original Fast-RCNN or Faster-RCNN code in Caffe or some other high-quality implementation.

[Reply](#)

i want to concatenate the pooling result of multiple convolutional feature maps to generate the final pooling features for detection tasks. Specifically, features from multiple lower-level convolution layers are ROI-pooled and those resulting features are then concatenated. can i do this ?

[Reply](#)

## Leave a Reply

Want to join the discussion?  
Feel free to contribute!

Name \*

Email \*

☐ I agree to allow deepsense.ai sp. z o.o. to process my personal data for the communication purposes relating to the services offered by deepsense.ai.

You can modify your privacy settings and unsubscribe from our lists at any time (see our [privacy policy](#)).

## SERVICES

Start with AI

AI solutions

AI services

Team augmentation

Machine learning

training

## PRODUCTS

Neptune

Seahorse

## ABOUT US

Blog

Company

Management

Scientific advisory

board

Careers

Press center

## CONTACT

contact@deepsense.a

